

Telemetry Blindness at the Agent–Business Boundary: A Controlled Measurement of Output-Contract Enforcement

Estefano Senhor Ferreira

Senior Software Engineer

estefano.senhor@gmail.com · <https://github.com/estefano-ferreira/authored-plan-agents>

Draft — extracted from repository commit `ae9e0f0`, 2026-08-07

Abstract

LLM agents that write business state can report success while persisting corrupted records. We reproduce this failure as a controlled ten-cell matrix (plus a control arm) — *decoding constraint* (provider-enforced response schema vs. schema stripped) crossed with *fault source* (an injector forcing format violations; the model’s natural behavior) and three *boundary repair policies* (an absorbing fallback, a genuine fence-stripping extractor, and a strict retry-once-then-fail output-contract guard) — accumulated as pre-registered cells over an originally non-factorial four-cell comparison, on one codebase, with the main comparison on one model; the grid is complete for the absorbing and strict policies. Pre-registered arms additionally replicate the decisive cells across two further model configurations, a second system-of-record engine, and a boundary strengthened to full-schema validation — the blindness signature is unchanged throughout — and the natural-violation cells replicate the originating incident under tracked evidence (natural rate 10/10, counters at zero). Ground truth is read directly from the persisted database of an external system of record. Two findings carry the paper. First, *telemetry blindness*: with repair positioned before the contract check, the tolerant configuration reported success 10/10 and persisted ten garbage business records while every platform-side instrument — status, contract counters, repair count — read success or no violation; only the database showed the corruption. The signature reproduced unchanged with the schema active and the fault injected downstream of it — a pre-registered follow-up reported here as a verification, its registered contrary expectation having rested on a misreading of the instrument (the fallback absorbs, never strips), corrected by a dated erratum in the pre-registration. Second, the guard’s price sheet: under enforced decoding both boundary policies were byte-equivalent and the guard’s idle cost was zero (natural violations 0/10 — a bounded observation, not an elimination claim); when a violation does arrive, the strict guard’s price is one extra model call, buying typed, unpersisted failure instead of silent corruption — a trade a naturally occurring provider schema-dialect mismatch exercised for us. The implication is deliberately narrow: constrained decoding is the suppression mechanism; the architectural boundary is what keeps the channel closed when that simpler mechanism silently breaks — and, measured against a competent repairer rather than an absorbing one, what it uniquely buys is typed visibility, not recovery superiority. All evidence is version-controlled and inspectable.

1 Introduction

The failure this paper measures was first observed in the wild, not designed. Ten consecutive executions of an email-to-ERP business flow reported success; ten times, the ERP received a truncated, markdown-fenced JSON blob as the customer-facing `summary` field — indistinguishable from a legitimate record to any downstream consumer. The platform’s contract counters read zero violations: a tolerant parser “repaired” each response before any check could see it. Every authorization held. The business API’s own validation passed — the corruption was in *content*, and content of valid form.

The gap has a name older than agents. Transaction processing drew this line four decades ago: in the paper that coined ACID, a committed transaction “by definition commits only legal results” — and, in the same discussion, “since, by definition, each transaction is correct, the effects of an inevitable incorrect transaction (i.e., the transaction containing faulty data) can only be removed by countertransactions” [1]. Consistency guarantees legality by the database’s lights; correctness of the written content is *assumed*, not guaranteed. What has changed is the writer, not the guarantee. Under programmatic writes the gap is rarely exercised: code writes what the programmer specified, so a semantically wrong value of valid form is the exception. Under LLM generation, form is enforced at the decoder while content is probabilistic — the gap goes from rarely exercised to routinely exercised. That is why business validity, which could remain an implicit programmer responsibility, needs an explicit owner when the content author is stochastic.

The phenomenon is not idiosyncratic. Advani measures *false success* — an agent claiming completion with no corresponding state change in the system of record — at 45–48% of failures in single-control τ^2 -bench domains, and finds LLM judges fail to detect it (AUROC ≤ 0.65) [2]. Cao et al. find 27–78% of benchmark-reported successes to be procedurally corrupt [3]. SABER localizes the damage: deviations in *mutating* steps reduce task-success odds by up to 92–96%, while deviations in non-mutating steps have little effect [4] — the write path is where failures become decisive.

What the literature does not settle is an engineering question. Provider-enforced structured output (constrained decoding) now exists and demonstrably works — so the question that can actually be answered with measurements is conditional: *when decoding enforcement silently fails, what does each boundary policy between generated content and the business write do with the violation — and what does keeping the stricter policy cost while decoding still works?* This paper answers with a controlled measurement whose ground truth is not the agent’s report, not the platform’s telemetry, but rows in the business database.

Contributions. (1) A controlled comparison of ten boundary configurations — decoding enforcement crossed with fault source (injected, natural) and three boundary repair policies (absorbing, strict, extracting), plus a shared control arm, accumulated as pre-registered cells over an originally non-factorial four-cell design — with ground truth read from a persisted external system of record. The grid is complete for the absorbing and strict policies; the extractor column carries injected-fault cells only (§3). State-grounded scoring itself is established practice — deterministic evaluation against the persisted environment is how STATE-Bench scores state-mutating tasks [8] and what GroundEval provides as an LLM-judge replacement [9]; what this comparison adds is the *question* asked of that ground truth: not whether the agent reached the correct final state, but what each boundary policy did to the *content* of the records it persisted — a mechanism ablation reading record quality row by row, priced on both sides. (2) The measured finding that *tolerant repair conceals rather than repairs*, and specifically that it blinds the platform’s own integrity counters — the audit trail records zero violations in the cell that persists 100% garbage: the *telemetry blindness* defined in §4. (3) An honest pricing of the boundary guard as insurance: idle cost zero when decoding enforcement works, one extra model call per violation when it does not, and a naturally occurring configuration failure in which the insurance paid out.

2 The measured system

The reference implementation is a minimal agent platform for business process automation (three agents, eight tools, two human-authored plans) in which non-determinism is confined to three model decisions: plan selection, capability selection, and content generation. Step order is

declared in versioned YAML; the model never decides what runs next. Three properties matter for this experiment:

1. **External system of record.** Business writes go through an ERP service (SQLAlchemy over SQLite) that owns its invariants as database constraints and rejects violations with typed errors. The platform holds no business entity. Ground truth for every cell is read from this database, never from platform state.
2. **Output contract on the write path.** The *boundary* is the single interface through which generated content becomes business state. The declared JSON Schema is enforced at the *decoding* level (the provider-side `responseSchema`); for the main comparison the boundary itself applies a strict decode-shape check — `json.loads`, a bare-JSON-object test, and presence of the contract’s required keys, with no type, enum, format, or `additionalProperties` validation — to the content that feeds a business write, at a point the generator cannot bypass; on violation, at most one retry, then a typed failure — nothing unvalidated persists. A pre-registered follow-up arm (§6) then strengthened the boundary to full validation of the declared schema — exact key set, string types, enum — and re-ran the schema cells, so the positional result is now measured at both validation strengths. This is the mechanism under test.
3. **Unsampled audit.** Proposals, policy decisions, business validations, output-contract events and compensations are recorded append-only, as distinct kinds. The experiment measures, among other things, whether this trail can be trusted — it cannot, in one cell, for a reason worth stating precisely (§4).

The pattern’s full specification and its seven load-, registration- and runtime-enforced restrictions are in the repository; none of them is the subject here. The subject is the boundary between generated content and persisted business truth.

3 Experimental design

Axes. *Decoding constraint:* the provider-enforced JSON schema (`responseSchema`) either passed through to the model provider, or stripped by an infrastructure decorator (`SchemaStrippingClient`). *Boundary policy:* either a strict output-contract guard (`retry_once_then_fail`) or a tolerant repair decorator (`TolerantRepairClient`) implementing an *absorbing fallback*: for any generation response that does not parse as a bare JSON object, it discards nothing and strips nothing — it stuffs the original text, fences included, into the contract’s fields (`summary` set to the first 200 characters, `reply_body` to the full text), producing a JSON object the downstream parser always accepts and passing the result through — an emulation, as an infrastructure decorator, of the historical fallback behavior that produced the original incident. Construct validity requires one precision: the historical fallback lived *inside the capability’s parser* (after a parse failure it stuffed the fenced blob into the business field); the emulation reconstructs that behavior as an infrastructure decorator so that measurement never touches the platform’s behavior layer. The comparison is therefore against a reconstruction at a different layer — and what transfers is positional, not layer-specific: any absorption upstream of the contract check blinds the counters the same way (Figure 2).

Fault injection. Cells A and B (no schema) additionally run a generation-fault injector (`GenerationFaultInjectionClient`) that re-fences *every* generation response. The violation rate in these cells is therefore forced by construction — 10/10 measures the injector, not the model. This is deliberate: the natural no-schema violation rate on this model was measured in an earlier round (10/10 fenced — itself an artifact of flash-tier defaults under prompt-only

control), and the question this matrix asks is conditional: *given* a violation, what does each configuration do with it? The injector also re-fences the retry attempt, so recovery in cell B is impossible by construction; the guard’s failure path, not its recovery path, is what B measures.

Stated plainly: because the injector originally ran only in the no-schema cells, **the original four-cell design was not a complete factorial** — the fault arm co-varied with the decoding axis. Cells C/D measured the schema configuration under *natural* conditions only; within the original four cells, no configuration subjected active decoding enforcement to a forced violation. Successive pre-registered cells have since closed that gap: E/F injected the fault *downstream* of the active schema, and G/H repeated that injection under a genuine extractor rather than the absorbing fallback; I/J then measured the no-schema arm’s *natural*, uninjected violation rate under each boundary policy (all in §6) — together completing the decoding-constraint × fault-source grid for the absorbing and strict boundary policies. The extractor column remains populated only under injected faults; the findings are worded accordingly (F3 claims invisibility of the repair axis and zero idle cost, not that decoding “holds under attack”).

A third boundary policy: the genuine extractor. Two further pre-registered cells (G and H; docs/preregistration-extractor-repair-cells.md, committed before the instrument existed) added a genuine fence-stripping extractor, `ExtractorRepairClient`, as a third value on the boundary-policy axis alongside the absorbing fallback and the strict guard: for a generation response matching a single markdown fence envelope (optional language tag), one pass, no recursion, it replaces the content with the captured inner text and records one repair; anything that does not match — bare JSON, prose, a double-fenced envelope’s outer layer once stripped — passes through unchanged, with no absorption and no field-stuffing. Cells I and J (docs/preregistration-natural-violation-cells.md, committed before the runner change) close the design’s other remaining gap: the no-schema arm’s natural, uninjected violation rate against the absorbing and strict policies. Results for both pairs are in §6.

Cells and control. Ten repetitions per cell of the two-agent email-to-ERP plan (the only plan with a generation step), each against a fresh SQLite file. A control arm runs the D configuration on the no-generation scheduling plan (3 repetitions): the axes should touch nothing but the generation step. Model: `gemini-3.1-flash-lite`, same prompts throughout; everything varies by command-line flags only.

What is read as the result. Reported status; contract violation/retry/failure counters from the audit trail; repair count; rows created in the ERP database; per-row validity of the persisted **summary** (checked by direct SQL inspection, not via any platform report); tokens and cost.

Price basis. Cost figures are computed by the runner from the per-model pricing table in the platform’s observability writer: for `gemini-3.1-flash-lite`, \$0.25 per million input tokens and \$1.50 per million output tokens, rates as recorded on 2026-07-28 from the official pricing page (<https://ai.google.dev/gemini-api/docs/pricing>). Token counts come from the provider’s own `usage_metadata` (`prompt_token_count`, `candidates_token_count` plus thought tokens where the model reports them, `cached_content_token_count`), not from an estimator. Cached input is billed at 10% of the input rate in the runner’s formula; in these runs cached tokens were zero — verified by recomputing all 243 recorded per-execution costs (the ten-cell matrix and its control arm, and the follow-up arms at their own list rates: `gpt-4o-mini` and Groq-served `gpt-oss-120b`, both \$0.15 / \$0.60 per million as recorded; the Postgres arm, the full-schema arms, and the extractor and natural-violation cells G/H/I/J all on the baseline model’s rates) from their token counts alone, which reproduces every recorded figure exactly (`scripts/verify_costs.py` in the repository). All runs executed inside free-tier quota, so the

Cell	Configuration	Status (10 reps)	Viol./Retry/Fail	Repairs	Rows	Valid	Cost
A	no schema + fault + tolerant	completed 10/10	0/0/0	10	10	0	\$0.00304
B	no schema + fault + strict	failed_clean 10/10	10/10/10	0	0	—	\$0.00509
C	schema + tolerant	completed 10/10	0/0/0	0	10	10	\$0.00247
D	schema + strict	completed 10/10	0/0/0	0	10	10	\$0.00246
ctrl	D config, no-generation plan	completed 3/3	0/0/0	0	3	n/a	\$0.00017
E	schema + fault + tolerant	completed 10/10	0/0/0	10	10	0	\$0.00245
F	schema + fault + strict	failed_clean 10/10	10/10/10	0	0	—	\$0.00384
G	schema + fault + extractor	completed 10/10	0/0/0	10	10	10	\$0.00248
H	no schema + fault + extractor	failed_clean 8/10, completed 2	9/9/8	19	2	2	\$0.00491
I	no schema, natural, tolerant	completed 10/10	0/0/0	10	10	0	\$0.00295
J	no schema, natural, strict	failed_clean 9/10, completed 1	10/10/9	0	1	1	\$0.00500

Table 1: The integrity matrix. “Rows” are records created in the external ERP database; “Valid” counts rows whose persisted **summary** is clean one-line text (verified by direct SQL inspection). Cell A’s ten rows each carry a double-fenced JSON blob in the business field while every counter the platform exposes reads zero. Below the first rule: the pre-registered follow-up cells E and F (§6), run after the original comparison — E reproduces A’s signature with the decoding constraint active upstream. Below the second rule: the genuine-extractor pair G/H and the natural-violation pair I/J (both §6) — G recovers valid content under the same injected fence that corrupted E, with the counters just as blind; H’s eight typed refusals are a double-fenced envelope defeating a single-pass extractor; cell I’s ten rows carry single-fenced JSON blobs produced by purely *natural* violations, no injector anywhere. Model calls per repetition are structural rather than tabulated: every **completed** repetition makes the plan’s fixed selection calls plus one generation call, and the strict cells’ **failed_clean** repetitions add exactly one retry generation call — the entire source of F4’s ratios. Latency means are omitted: they are contaminated by free-tier 429-backoff stalls (worst case 111s) and carry no signal about the axes.

figures are list-price equivalents, not amounts billed. List prices change; the figures are reproducible against the rates as of the stated date, and the token counts — which do not change — are in the tracked evidence files (`results/integrity-matrix/integrity_matrix.jsonl`, per (cell, rep)).

4 Results

F1 — without the schema, the boundary alone separates the outcomes. Same fault, same model, same pipeline: A reports **completed** 10/10 and persists 10 garbage business records; B reports **failed_clean** 10/10 and persists zero. Both outcomes are deterministic by construction, uniformly: the injector re-fences every response, B’s guard is coded to fail typed on it, and A’s absorbing fallback is coded to stuff the fenced text into the contract’s fields regardless of content. B *verifies*, under real-model conditions, that the strict policy does what its code says; A *verifies* the same for the tolerant policy — both are designed contrasts, not discoveries. What A contributes is not an unprogrammed behavior but an end-to-end demonstration, on a real pipeline: the corrupted content reaches the business database while every platform-side instrument reads clean (F2).

F2 — tolerant repair conceals, and blinds the telemetry. Cell A’s contract counters read 0/0/0 because the repair runs *before* the contract check ever sees the response; only the repair count (10) and direct database inspection reveal what happened. The practitioner literature warns in prose that syntactic repair “fixes syntax, not semantics”; part of it celebrates high repair-success rates as a virtue. The matrix quantifies the inverted reading: “repaired” 10/10, valid 0/10. The consequence extends a lesson the false-success literature draws about

	Cell A (no schema, tolerant)	Cell B (no schema, strict)	Cell E (schema, tolerant)
<i>what the platform reports</i>			
Agent report	completed 10/10	failed_clean 10/10	completed 10/10
Contract counters (viol. / retry / fail)	0 / 0 / 0	10 / 10 / 10	0 / 0 / 0
Repair count	10	0	10
<i>what the database holds</i>			
Rows in ERP	10	0	10
Valid rows	0	—	0

Figure 1: What each observer sees (data from Table 1; no new measurement). The separation between the two bands is the finding: in cells A and E, three platform-side instruments agree the runs succeeded — status **completed**, contract counters at zero, ten “successful” repairs — while the business database holds ten corrupted records and zero valid ones. Cells A and E differ in decoding condition (schema stripped vs. active, the fault injected downstream of it): the same blindness signature in both columns is Proposition 1’s positional claim (§4), visible at a glance.




Figure 2: Why F2 happens: one position in the sequence, not a property of repair. In the tolerant pipeline the repair runs before the contract check, so the check — and the counters it feeds — never see the violation; in the strict pipeline the check runs first and the violation is counted before retry or typed failure.

agent self-reports [2, 3] one level deeper: in a repair-before-check pipeline, *the platform’s own integrity telemetry is not evidence either*. Only the system of record is (Figure 1); and the mechanism is positional, not a property of repair itself (Figure 2). Two levels of this claim must be kept apart. The *positional* mechanism is general: any absorption upstream of a check blinds the counters that check feeds. The *measured* result is narrower: absorption of format violations ahead of a format check. A repair that genuinely normalizes content ahead of a form check, or form repair ahead of a semantic check, was not tested and may behave differently (§6).

Definition 1 (Telemetry blindness). A pipeline exhibits *telemetry blindness* when every platform-side observer reports success or no violation while the system of record holds corrupted state.

Proposition 1 (Sufficient condition). Given a violation that reaches persistence, any absorption step positioned upstream of the check that feeds those observers, *and whose own telemetry reports its operation as success*, suffices: the check never sees the violation, so the counters it feeds cannot record it — and the absorption step’s own counter reads as successes,

not violations. The success-reporting clause is not incidental: an absorption step that classified its own interventions as violations would feed an observer that reads unhealthy, breaking Definition 1’s premise by construction — but success-reporting is precisely what repair semantics mean, and every repairer measured here carries it. With that scope, the condition is mechanism-agnostic: it holds for any such absorption step, not only the absorbing fallback measured here (the measured narrowing above stands).

Observed instance. Cell A: agent report `completed`, contract counters at zero, repair count reporting ten successes, ten corrupted rows in the business database. A second instance was measured after this definition was written and pre-registered (§6): the same signature — counters at zero, ten repairs, ten corrupted rows — with decoding enforcement active upstream and the fault injected downstream of it. A third instance was measured under a further pre-registered arm (§6): the same signature again, with the boundary performing full validation of the declared schema — the absorbed payload is schema-valid by construction, so no boundary-side validation strength can reject it; this outcome was determined by the code before the run and is reported as a verification. A fourth instance was measured under the natural-violation cells (§6): cell I reproduces the same signature — counters at zero, ten repairs, ten corrupted rows — with *no injector anywhere in the pipeline*; the violations are entirely natural, making cell I the first instance this study produced without a fault injector. The condition is positional, indifferent to what sits upstream of the absorption, to how strong the downstream check is, and to whether the violation was forced or occurred on its own. Instance counting, fixed here: an *instance* is a distinct (decoding condition, boundary strength) configuration exhibiting the signature — four so far; engine replications of these instances (§6) are counted as replications, not further instances.

F3 — with the schema, no violation occurred and the repair axis is invisible. Cells C and D are byte-equivalent in outcome (10 clean rows each, $\sim$ same tokens, $\sim$ same cost) with zero natural violations at $N=10$ — an observation bounded at $\leq 25.9\%$ (§6), not an elimination claim, and these cells never faced a forced violation (§3). The direct evidence that decoding-level enforcement suppresses *natural* violations is the earlier before/after round on the same pipeline: 10/10 violations under prompt-only control $\rightarrow$ 0/10 with the schema, at -18% cost per completed task. Cells I and J (§6) now measure the same no-schema natural rate under tracked evidence at this date — 10/10 and 10/10 (20/20 first attempts) — replacing reliance on that earlier round’s surviving aggregates with a directly inspectable record. What C/D add is the guard’s price sheet: under the enforced-decoding configuration the guard has literally nothing to do, and its idle cost — extra model calls, extra tokens — is zero. The pre-registered full-schema arm (§6) re-measured this idle cost at full validation strength — still zero: cells C/D remain byte-equivalent under the strengthened check, with no additional model calls.

F4 — the guard’s price is one extra call per violation. B spent $1.92\times$ A’s output tokens (2636 vs. 1376) and $1.68\times$ its cost (\$0.00509 vs. \$0.00304 per 10 reps): the retry doubles the generation call and both attempts fault by construction. That is the entire price of knowing. The natural-violation cells reproduce the same ratio without any injector: J spent $1.69\times$ I’s cost (\$0.00500 vs. \$0.00295 per 10 reps).

F5 — the control arm isolates the generation step. The no-generation plan completed 3/3 with clean persisted state under the D configuration: the axes touch nothing else in the pipeline.

5 The naturally occurring case

The matrix’s guarded no-schema cell models a configuration class that occurred naturally during the study, unplanned. The first corrected measurement round passed the output schema through the provider’s `response_schema` config field — which accepts an OpenAPI-subset dialect and rejected the standard JSON-Schema key `additionalProperties` with a live HTTP 400. Constrained decoding was thereby silently disabled by configuration: exactly the class of provider-side surprise that had previously produced 10/10 silent corruptions. With the boundary guard in place, it produced instead 10/10 *visible, typed* failures with the ERP at zero rows. The fix (the provider’s `response_json_schema` field, which accepts real JSON Schema) is one line — and the secondary finding is portable: JSON Schema is not portable verbatim across providers (two Gemini config fields with different grammars; Anthropic requiring `additionalProperties: false` plus `required`; OpenAI’s `strict` wrapper), and dialect translation is adapter responsibility.

This is the paper’s claim in one incident, stated narrowly: the architectural boundary is *not* the primary suppression mechanism — decoding-level enforcement is what suppressed every natural violation we measured (F3 and the before/after round). The boundary is what keeps the channel closed when the simpler mechanism silently breaks. As insurance it is unusually well priced: idle cost zero (F3), one extra model call per violation when it fires (F4). We observed one naturally occurring claim event; its base rate in production is unknown ($n=1$), and we state it as an existence proof, not a frequency.

6 Limitations and statistical bounds

Small cells, stated as such. Ten repetitions per cell. A 0/10 observation bounds the underlying rate only to $\leq 25.9\%$ at one-sided 95% confidence ($1 - 0.05^{1/10}$); “zero violations at $N=10$ ” is compatible with a material residual rate and is reported as an observation, not an elimination claim. The qualitative cell separations (10 garbage rows vs. zero; counters blind vs. counting) do not depend on rate estimates.

The injector measures the injector. The 10/10 violation rate in A/B is forced; the natural rate on this model under prompt-only control (10/10 in the original round) is an artifact of flash-tier defaults. All conditional claims are of the form “given a violation”.

Not a complete factorial — and what the completed arm measured. At the time of the four-cell comparison the fault arm existed only in the no-schema column (§3). The two missing cells — schema + injected fault, under each boundary policy — were subsequently pre-registered (branches and the uncomfortable outcome fixed before any number existed) and run on the same model, 10 repetitions each. The pre-registration recorded a contrary expectation: because the injector corrupts downstream of the provider, the injected fence wraps schema-valid JSON, and we predicted — in the pre-registration, in review, and in an earlier version of this paragraph — that fence-stripping repair would plausibly recover *valid* content there. The run’s numbers contradicted that expectation. A post-run reading of `TolerantRepairClient`’s source showed, however, that the outcome was determined by the instrument before any number existed: the client never strips fences at all; for any response that does not parse as a bare JSON object it absorbs the original text — fences included — into the contract’s fields. The pre-registration carries a dated erratum (2026-08-07, `docs/preregistration-schema-fault-cells.md`) recording this correction. The tolerant cell reported **completed** 10/10 and persisted ten fenced blobs into the business field with contract counters at zero and ten repairs reported — telemetry blindness reproduced under an *active* decoding constraint, a second observed instance of Proposition 1 in a different decoding condition; cell E is accordingly reported as a *verification*

that the historical fallback persists garbage independently of the decoding condition, not as a falsified prediction. The strict cell reported `failed_clean` 10/10 with zero rows, refusing content that was, under the fence, valid — and remains the verification it was pre-declared to be (the injector re-fences the retry, so its outcome is deterministic by construction). The counterfactual “tolerance would have recovered it” was therefore never open to begin with: it was settled by the instrument’s construction and confirmed, not falsified, by the run — the alternative to refusal was not recovery; it was corruption. The no-schema/no-fault arm — natural violations under each boundary policy, the design’s longest-standing admitted gap — has since run as cells I and J, reported in their own paragraph below (natural rate 10/10 in both cells); the one factorial gap that remains is the extractor policy under natural, uninjected violations, which has not run.

The full-schema boundary arm. The v0.4.0 review asked the strong form of the boundary question: does the positional result survive when the boundary performs full validation of the declared schema — exact key set, string types, the `request_type` enum — instead of the decode-shape check? The arm was pre-registered under a new discipline adopted the same day: a mandatory “what the code already determines” section, written before any prediction, so that a deterministic outcome cannot later be dressed up as a discovery (`docs/preregistration-boundary-schema-validation.md`). Determination: cells E and F were code-determined before the run — the absorbing fallback’s payload passes full schema validation by construction (it supplies exactly the three declared keys, all strings, with `request_type` in the enum), and the injector re-fences the retry, defeating the strict guard’s recovery path exactly as in the original F — and are reported here as verifications, not discoveries. The open cells were C and D: under active decoding, does the strengthened check reject content the shape check would have passed? The run confirmed both registered predictions. P1: zero full-schema-only rejections in ten natural repetitions per cell (0/10 each, bounding the underlying rate at $\leq 25.9\%$ one-sided 95%, as everywhere in this study), with C and D remaining byte-equivalent (10 clean rows each, $\sim$ same tokens, $\sim$ same cost). P2: the strengthened check’s idle cost is zero — no additional model calls in C/D relative to the recorded baseline C/D. Consequence stated plainly: “position, not strength” is no longer only a structural argument — it is measured at the strongest boundary check the declared schema admits. A further pre-registered engine replication (`docs/preregistration-fullschema-postgres.md`) then re-ran this arm’s four cells against PostgreSQL — its registered Branch A (full transfer) materialized exactly: C/D 10/10 clean, E ten fenced garbage rows with counters at zero, read from the live Postgres by direct SQL, F 10/10 typed failures with zero rows — closing a scope gap this paper had carried silently: cells C–F had never run against the second engine.

The genuine-extractor cells. The v0.5.0 external review (M1) observed that this study had measured only an *absorbing* repairer — the historical fallback that stuffs the raw response into the contract’s fields — while the practitioner literature’s default repair is *extraction* (markdown fence removal). Two cells, pre-registered under the same discipline (`docs/preregistration-extractor-repair-cells.md`, committed before the instrument existed), completed the repair-policy axis with a genuine extractor, `ExtractorRepairClient` (single-pass markdown fence removal, no absorption, no field-stuffing). Determination, fixed before any number: cell G is determined conditional on decoding conformity — under an active schema the provider returns schema-valid JSON (measured 0/10 natural violations at both validation strengths in C/D), the injector adds one fence, and a single-pass extractor matches it by construction, so the recovered content passes full validation; cell H is genuinely open, because without the schema the content under the fence is unconstrained and depends on whether the model itself self-fences. Results confirmed the determined branch and scored the open one. **Cell G: completed** 10/10, contract counters 0/0/0, 10 repairs, ten rows, all ten valid. **Cell H: failed_clean** 8/10 (2

Arm	Varies		Cells	Reading
<code>gpt-4o-mini</code>	model family		A/B	transfer on every boundary-reaching rep; 13/20 diverge upstream
<code>gpt-oss-120b</code>	family (open weights)		A/B	full transfer, full boundary samples
<code>-pg</code>	SOR engine		A/B	full transfer (replication)
<code>fullschema</code>	boundary strength		C–F	P1/P2 confirmed; blindness unchanged (verification)
<code>fullschema-pg</code>	engine, at full strength		C–F	full transfer (replication)
extractor cells	repair mechanism		G/H	G recovers valid content with counters blind (verification); H 8/10 refused, double fence defeats single pass
natural cells	fault source — none		I/J	natural rate 10/10; blindness instance four; one natural retry recovery

Table 2: The pre-registered arms and follow-up cells at a glance. “Reading” compresses the paragraphs above; every arm row’s evidence is the arm-suffixed file set in `results/integrity-matrix/` and its pre-registration under `docs/`; the extractor and natural rows are baseline-namespace cells (G/H/I/J, no evidence suffix), each with its own pre-registration. Instances vs. replications per the counting rule fixed in §4.

completed), contract counters 9/9/8, 19 repairs, two rows, both valid. The mechanics behind H are code-coherent: in 8 of 10 repetitions the model self-fenced its own response (implied self-fencing rate 9/10 on first attempts) on top of the injector’s fence, producing a double-fenced envelope a single-pass extractor cannot fully unwrap on either attempt (16 of H’s 19 repairs); one repetition recovered on retry (2 repairs) and one recovered on the first attempt (1 repair). The registered prediction (`failed_clean` $\geq 8/10$) held exactly at the threshold. The committed framing consequence (Branch A) is reported in the terms it was pre-registered in: the economics of refusal-versus-repair is *repairer-specific* — the absorbing fallback corrupted cell A/E’s content, the genuine extractor recovered cell G’s, and the contract counters were blind in both. What the strict guard buys is therefore not superiority over any repair; it is typed visibility. A correct extractor beats both the absorber and the guard on availability in cell G, while leaving the counters exactly as blind as the absorber does.

The natural-violation cells. The design’s longest-standing admitted gap — the no-schema/no-fault arm, the model’s natural violation behavior under each boundary policy, never run as a matrix cell — closed under the same pre-registration discipline (`docs/preregistration-natural-violation-cells.md`, committed before the runner change): cells I (no schema, no injector, tolerant) and J (no schema, no injector, strict), 10 repetitions each, no new instrument, only configuration. The natural violation rate measured 10/10 in both cells. **Cell I: completed** 10/10, contract counters 0/0/0, 10 repairs, ten rows, **zero** valid — the fourth observed instance of Proposition 1, and the first produced with no injector anywhere in the pipeline. **Cell J: failed_clean** 9/10 (1 **completed**), contract counters 10/10/9, zero repairs, one row, valid. Nine of J’s ten repetitions ended in a typed refusal after the retry also violated; the tenth recovered — the model returned bare JSON on the second, identical call. That nuance matters: under an injector the retry is doomed by construction (the injector re-fences every attempt, cells B/F/H); under a natural violation nothing forces the second attempt to repeat the first, and here, once, it did not. The combined first-attempt violation rate across I and J is 20/20, replicating the originating incident round’s 10/10 under tracked, inspectable evidence rather than the lost baseline’s surviving aggregates — the loss itself (below) stays disclosed, not retracted.

One violation class. The injected fault (markdown re-fencing) is the class the original incident exhibited. Tolerant repair in the wild absorbs other classes (truncation, wrong fields, schema-valid-but-wrong content); semantically wrong content of valid form passes every cell clean — content quality inside a valid schema is explicitly out of scope, and the matrix measures *form* integrity only.

The baseline’s raw file is lost. The before/after round cited in F3 rests on a baseline whose raw per-execution file was deleted by a cleanup step before the repository was under version control. Its aggregates survive in the repository’s study log, and the corrected round’s raw file is tracked; the 10/10 figure is therefore reproducible from surviving aggregates, not from raw records. The natural-violation cells I and J (above) now provide tracked, inspectable evidence for the same phenomenon at this date, so no current claim in this paper rests on the surviving aggregates alone. The repository discloses this loss, and so does this paper.

Model scope — one model in the main comparison, arms beyond it. All cells of the main comparison ran on `gemini-3.1-flash-lite`. A pre-registered cross-model replication of cells A/B has since run its family arm (`gpt-4o-mini`, 10 repetitions per cell, model-suffixed tracked evidence): every repetition that reached the generation boundary reproduced the original reading — tolerant: 6/6 completed with fenced garbage persisted and counters at zero (telemetry blindness in a third model family); strict: 1/1 typed failure — and zero rows persisted in all 10 strict repetitions regardless of failure locus. The caveats are part of the result: the strict cell’s boundary sample is one, because 13/20 repetitions diverged *upstream* on selection-response format (the model answers a sentence instead of the bare id) — a per-family selection result reported separately per the pre-registration, and itself a falsification of this study’s earlier reading that the id-only selection contract is “naturally robust” as a property of the task. A further pre-registered arm (Groq-served `gpt-oss-120b`, open-weights) then provided what the mini arm could not: *every* repetition reached the boundary in both cells (the id-only protocol held in this configuration), giving full boundary samples — tolerant: 10/10 garbage persisted with counters at zero; strict: 10/10 typed failures, zero rows. The mini arm’s sample-of-one caveat is thereby resolved by measurement in a different configuration, not withdrawn. The frontier-tier arm (`claude-sonnet-5`, amendment recorded) has not run.

Self-built system of record. The ERP is a mock written by the same author, with two invariants designed to be checkable. The pattern’s guarantee delegates wholesale to the system of record’s competence; this study cannot measure an incompetent one. A pre-registered engine arm reran cells A/B with the same ERP on PostgreSQL (baseline model held fixed) and reproduced both readings exactly, and a second pre-registered engine replication later re-ran the full-schema arm’s cells C–F on PostgreSQL with the same result (full transfer) — the results are engine-portable and a SQLite artifact is ruled out — but the circularity itself is untouched: the engine varied, the owner did not. Relatedly, the platform’s registration-time checks validate *coherence of declared* tool metadata, not the truth of the declarations.

7 Related work

False and corrupt success. Advani [2] names and measures false success and shows detection failing; Cao et al. [3] measure procedural corruption behind reported success; SABER [4] shows mutating steps are where deviations become decisive. This experiment is a persistence-grounded replication of the phenomenon with a mechanism ablation none of the three performs.

Structured output. Provider-enforced structured output and constrained-decoding libraries are commodity; the known enforcement-level gradient (prompt-only weakest, native structured

output strongest) is replicated, not discovered, by our before/after. What we add is the controlled comparison against the boundary policy and the database-grounded reading. The closest experimental neighbor is harness engineering for auditable enterprise agents [5]: contract preservation measured at a replaceable composition boundary, with a fault-injection control confirming the validators flag deliberately broken contracts (270 boundary runs across three hosted models) and an ablation in which code-owned enforcement preserves both safety and full utility where a bolt-on guardrail over-refuses (88/120 vs. 120/120). Their contracts govern the composed *answer* (source grounding, entity routing, output hygiene); ours governs a business *write*, with ground truth read from the persisted store and a repair-policy axis their design does not vary — the telemetry-blindness result (F2) has no counterpart there.

Transaction processing. The distinction this paper leans on is the transaction literature’s own: committed results are “legal” by the database’s lights, and the transaction containing faulty data sits outside the guarantee, removable only by countertransactions [1]. Cell A is that distinction exercised one layer up — constraints validated, transaction committed, content wrong. The analogy is claimed for the gap only: nothing in this paper provides atomicity, isolation, or durability.

State-grounded evaluation. Judging agents against persisted state rather than their own reports is established: STATE-Bench’s deterministic scorer “compares the final environment state to ground truth” for state-mutating tasks over a pre-populated database of business artifacts [8], and GroundEval is, by its own title, a deterministic replacement for LLM-as-judge in stateful agent evaluation [9]. Both ask whether the agent reached the correct state — task scoring. This paper asks a different question of the same ground truth: what each boundary policy did to the content of the records it persisted. The source of truth is shared inheritance; the mechanism ablation over record quality is what this measurement adds.

Gray failure and differential observability. The systems literature named the observability structure underneath our finding in 2017: a system experiences *gray failure* “when at least one app makes the observation that system is unhealthy, but observer observes that system is healthy” — *differential observability*, in a model where the observer is part of the system and the app is external to it [12]. Telemetry blindness is not that situation; it is its degenerate case. In cell A *no* entity observes unhealthy: the platform’s observer instruments (status, contract counters, repair count) read clean, and a downstream consumer of the ERP — the model’s app — reads records of valid form. The observability table of Huang et al. would place cell A in its no-failure quadrant; the disagreement is not between two observers but between every observer and the persisted state itself, and only direct inspection of the system of record reveals it. We therefore claim telemetry blindness as a specialization of differential observability — the case where the differential has collapsed — not as a new phenomenon. The structure is now actively benchmarked at the trace layer: AgentTelemetry maps fourteen fault types onto agent-specific span kinds and shows a complete span taxonomy lifting fault detection from 0.429 to 1.000 [10]. Telemetry blindness is the case span completeness does not reach: the absorbing step emits a *successful* span by its own contract, so no added span kind turns red — only the content of the persisted record does. At the agent layer, the nearest neighbour is a longitudinal taxonomy of silent failures in a production LLM agent runtime, itself framed in gray-failure terms and citing Huang et al., whose error-swallowing-and-dilution class the absorption step measured here instantiates [13]. Its “fail-plausible” class is a *different* specialization of the same lineage, and the distinction matters: there the failure fabricates narrative that deceives the observer; in telemetry blindness nothing is fabricated — the violation is absorbed upstream and every observer honestly reads success. Otherwise, the difference is method — taxonomy drawn from a production runtime there, controlled measurement of one mechanism here.

Error hiding. The absorption step is family-adjacent to the classic error-hiding (error-swallowing) anti-pattern — catching a failure and continuing without recording it [11] — and the practitioner literature already states that anti-pattern’s observational consequence on the availability axis: an API that answers 200 for failures keeps availability dashboards green. The distinction Proposition 1 rests on is one level deeper: in error hiding an error occurs and is swallowed; in the measured absorption *no error occurs at the absorbing step* — the repair succeeds by its own contract and its counter records successes. Telemetry blindness requires no swallowed exception anywhere in the pipeline; this is the same no-fabrication move drawn against fail-plausible above.

Durable execution. Temporal-style runtimes and the agent checkpointers built on them guarantee that an *execution* completes, replays, and survives crashes; they say nothing about whether the persisted *content* is business-true. A durable workflow persists garbage perfectly durably and perfectly auditably — cell A is exactly that artifact. Validity authority is orthogonal to durable execution, not redundant with it.

Adjacent enforcement loci. Intent-governed authorization narrows tool manifests server-side and checks payload effects before execution [6] — the guarantee-belongs-to-the-validator argument on the authorization axis, where ours is on the validity axis. Proof-of-execution frameworks formalize “null effect on deny” and tamper-evident execution records under cryptographic assumptions [7]; our audit trail makes no such claim (F2 is precisely a measurement of how far platform-side records can be trusted — and where they cannot).

Delimitation. Statements of absence in this paper’s related-work mapping (“the telemetry-blindness result (F2) has no counterpart there”) are search- and time-bound as of 2026-08-06, over a literature that does not index by these axes; during this study’s own review rounds, an adjacent paper that predated one of our searches was found only by a later, independent pass. Two further conditions are declared at the same level. References [8] and [11] function as pointers, not evidence — a vendor blog post quoted at sentence level and an encyclopedic anti-pattern entry — and no absence claim in this paper rests on either. And this study’s measurements, analysis and internal review are the work of a single author; the errors its 2026-08-07 erratum corrects were surfaced by review external to the project, not by its internal process — the determination rule adopted that day (§6) is the structural mitigation, not a substitute for independent replication. Absence findings mean “not found by a stated search at a stated date”, never “does not exist”.

8 Conclusion

Between an LLM’s output and a business database there is a boundary, and this paper measured what each boundary policy does when a contract violation reaches it — forced by an injector in the unconstrained cells, naturally occurring in the incident and the dialect accident that bracket them. In every round we measured, decoding-level enforcement suppressed all natural violations — a claim now measured on both sides of the contrast: with the schema stripped and no injector, cells I and J found the natural violation rate at 10/10, the same rate the schema suppresses to 0/10 wherever it is active — and a strict architectural guard behind it costs nothing while that remains true. What the guard buys is the difference between the two remaining worlds: when the decoding constraint silently breaks — a dialect mismatch, a misconfiguration, a provider change — a tolerant pipeline persists garbage while reporting success and zeroing its own counters — the telemetry blindness of §4 — while a strict boundary fails loudly, typed, with nothing persisted. In the combination we measured — absorption ahead of a format check — repair-before-check did not repair; it concealed. With a genuine extractor

in the same position (cell G), repair-before-check *did* repair — and the counters stayed equally blind: the concealment is positional, the recovery is repairer-specific, and what the strict guard uniquely buys is typed visibility, not superiority over every possible repair. What generalizes is the position, not the repair: any absorption upstream of a check blinds the counters that check feeds (Figure 2), and that positional lesson is what travels beyond this stack. The evidence for every sentence above is a version-controlled artifact: per-cell database snapshots, append-only audit trails, and raw model-call logs, inspectable from a fresh clone.

Implications for platform design. Three design consequences follow from the measurements, stated as requirements the findings support rather than as a new architecture, at three different levels of evidential support. First, an output-contract guard on the business write path — measured here at decode-shape strength and, in the follow-up arm, at full-schema strength — is an architectural requirement, not an optimization: its idle cost was measured at zero in the tested configuration (a bounded observation, F3/F4, §6), and was re-measured at zero under the full-schema boundary arm’s strengthened validation (§6), and its per-violation price is one model call, so omitting it saves nothing measurable and forfeits the only loud-failure path this study observed. This requirement rests on two supports of different kinds: the measured price sheet (F3/F4) prices it, and the naturally occurring dialect accident (§5) — the one observed event in which the guard caught what the decoding constraint had silently stopped catching — is what argues it is *necessary*, an existence proof at $n=1$, stated as such. Second, platform-side telemetry must not be treated as integrity evidence — audit has to read the system of record, because the configuration in which every instrument agrees and is wrong is demonstrated, named, and cheap to produce — this requirement rests on the four demonstrated instances: cell A, cell E, cell E re-run under full-schema boundary validation, and cell I under purely natural violations with no injector anywhere (§6). Third, repair-before-check must be made structurally impossible on business writes — the mechanism is positional, so no policy applied after the absorption can recover the signal — this requirement rests on Proposition 1’s structural argument, of which cells A, E and the full-schema arm’s E are observed instances rather than its proof. The reference implementation’s pattern (Authored-Plan Agents) encodes these three as registration- and load-time restrictions with their own violation probes; its specification and test suite are in the repository for any platform that wants to adopt the requirements without adopting the implementation.

Data availability. Repository tag v0.6.0 carries the complete evidence; the exact source state of this PDF is the commit named on the title page (its parent by construction — a PDF cannot embed the hash of the commit that carries it). The tag provides: implementation, restriction-enforcing test suite (35 passed, 7 strict-xfail anti-vacuity probes — the extractor unit tests for `ExtractorRepairClient` account for the growth since v0.5.0’s 27), and tracked evidence under `results/integrity-matrix/` (checkpointed JSONL, consolidated report, per-cell SQLite snapshots `matrix-{A,B,C,D,control,E,F}.sqlite` plus the model- and engine-suffixed arm snapshots (`-gpt-4o-mini`, `-openai-gpt-oss-120b`, `-pg`), audit trails), the genuine-extractor and natural-violation cells’ own snapshots (`matrix-{G,H,I,J}.sqlite` and their audit trails, all in the baseline namespace, appended to the same `integrity_matrix.jsonl` and consolidated report; evidence of the pre-registrations in `docs/preregistration-extractor-repair-cells.md` and `docs/preregistration-natural-violation-cells.md`), the full-schema boundary arm’s own snapshots (`matrix-{C,D,E,F}-fullschema.sqlite` and the `-fullschema` JSONL, report and audit trails, evidence of the arm pre-registered in `docs/preregistration-boundary-schema-validation.md`; its PostgreSQL replication beside them with the layered `-fullschema-pg` suffix, per `docs/preregistration-fullschema-postgres.md`) and `results/runs/` (the naturally occurring dialect-failure round, `run-20260727-215429.json`; the corrected round, `run-20260727-215811.json`). The Postgres arm’s ground truth was

read from the live PostgreSQL database by direct SQL and exported as the tracked -pg SQLite snapshots noted above — the uniform format the reader scripts accept (see `docs/preregistration-postgres-sor.md`). One caveat carried from §6: the baseline round’s raw file (`run-20260727-212338.json`) does not exist — lost to a pre-versioning cleanup; its aggregates survive in the repository’s study log.

References

- [1] T. Härder and A. Reuter. Principles of transaction-oriented database recovery. *ACM Computing Surveys*, 15(4):287–317, 1983.
- [2] L. Advani. From confident closing to silent failure: Characterizing false success in LLM agents. arXiv:2606.09863, 2026.
- [3] Y. Cao et al. Beyond task completion: Revealing corrupt success in LLM agents through procedure-aware evaluation. arXiv:2603.03116, 2026.
- [4] A. Cuadron, P. Yu, Y. Liu, and A. Gupta. SABER: Small actions, big errors — safeguarding mutating steps in LLM agents. arXiv:2512.07850, 2025. Under review, ICLR 2026.
- [5] J. Ahn and M. Kim. From prompts to contracts: Harness engineering for auditable enterprise LLM agents. arXiv:2607.08028, 2026.
- [6] G. Zhu and C. Wang. Intent-governed tool authorization for AI agents. arXiv:2606.22916, 2026.
- [7] J. Rhodes and G. Kang. Proof of execution: Runtime verification for governed AI agent actions. arXiv:2607.05397, 2026.
- [8] Microsoft. Introducing STATE-Bench: A benchmark for AI agent memory. Microsoft Open Source Blog, 2026-05-19. Accessed 2026-08-06.
- [9] GroundEval: A deterministic replacement for LLM-as-judge in stateful agent evaluation. arXiv:2606.22737, 2026.
- [10] K. C. Balusu. AgentTelemetry: A fault detection benchmark and toolkit for LLM agent observability. In *Proceedings of the 3rd ACM International Conference on AI-Powered Software*, 2026.
- [11] Error hiding. Wikipedia. Accessed 2026-08-06.
- [12] P. Huang, C. Guo, L. Zhou, J. R. Lorch, Y. Dang, M. Chintalapati, and R. Yao. Gray failure: The Achilles’ heel of cloud-scale systems. In *Proceedings of HotOS*, 2017.
- [13] W. Wu. When errors become narratives: A longitudinal taxonomy of silent failures in a production LLM agent runtime. arXiv:2606.14589, 2026.